

Pintos threads and scheduling

Recap

- So far we know how to
 - run and debug pintos tests
 - add our own tests
 - make use of `lib/kernel/list.h`

Threads


- Baseline kernel (`pintos/threads/`)
 - Multiple threads running in kernel mode
 - Threads are similar to processes
 - Share the same address space

Thread control block

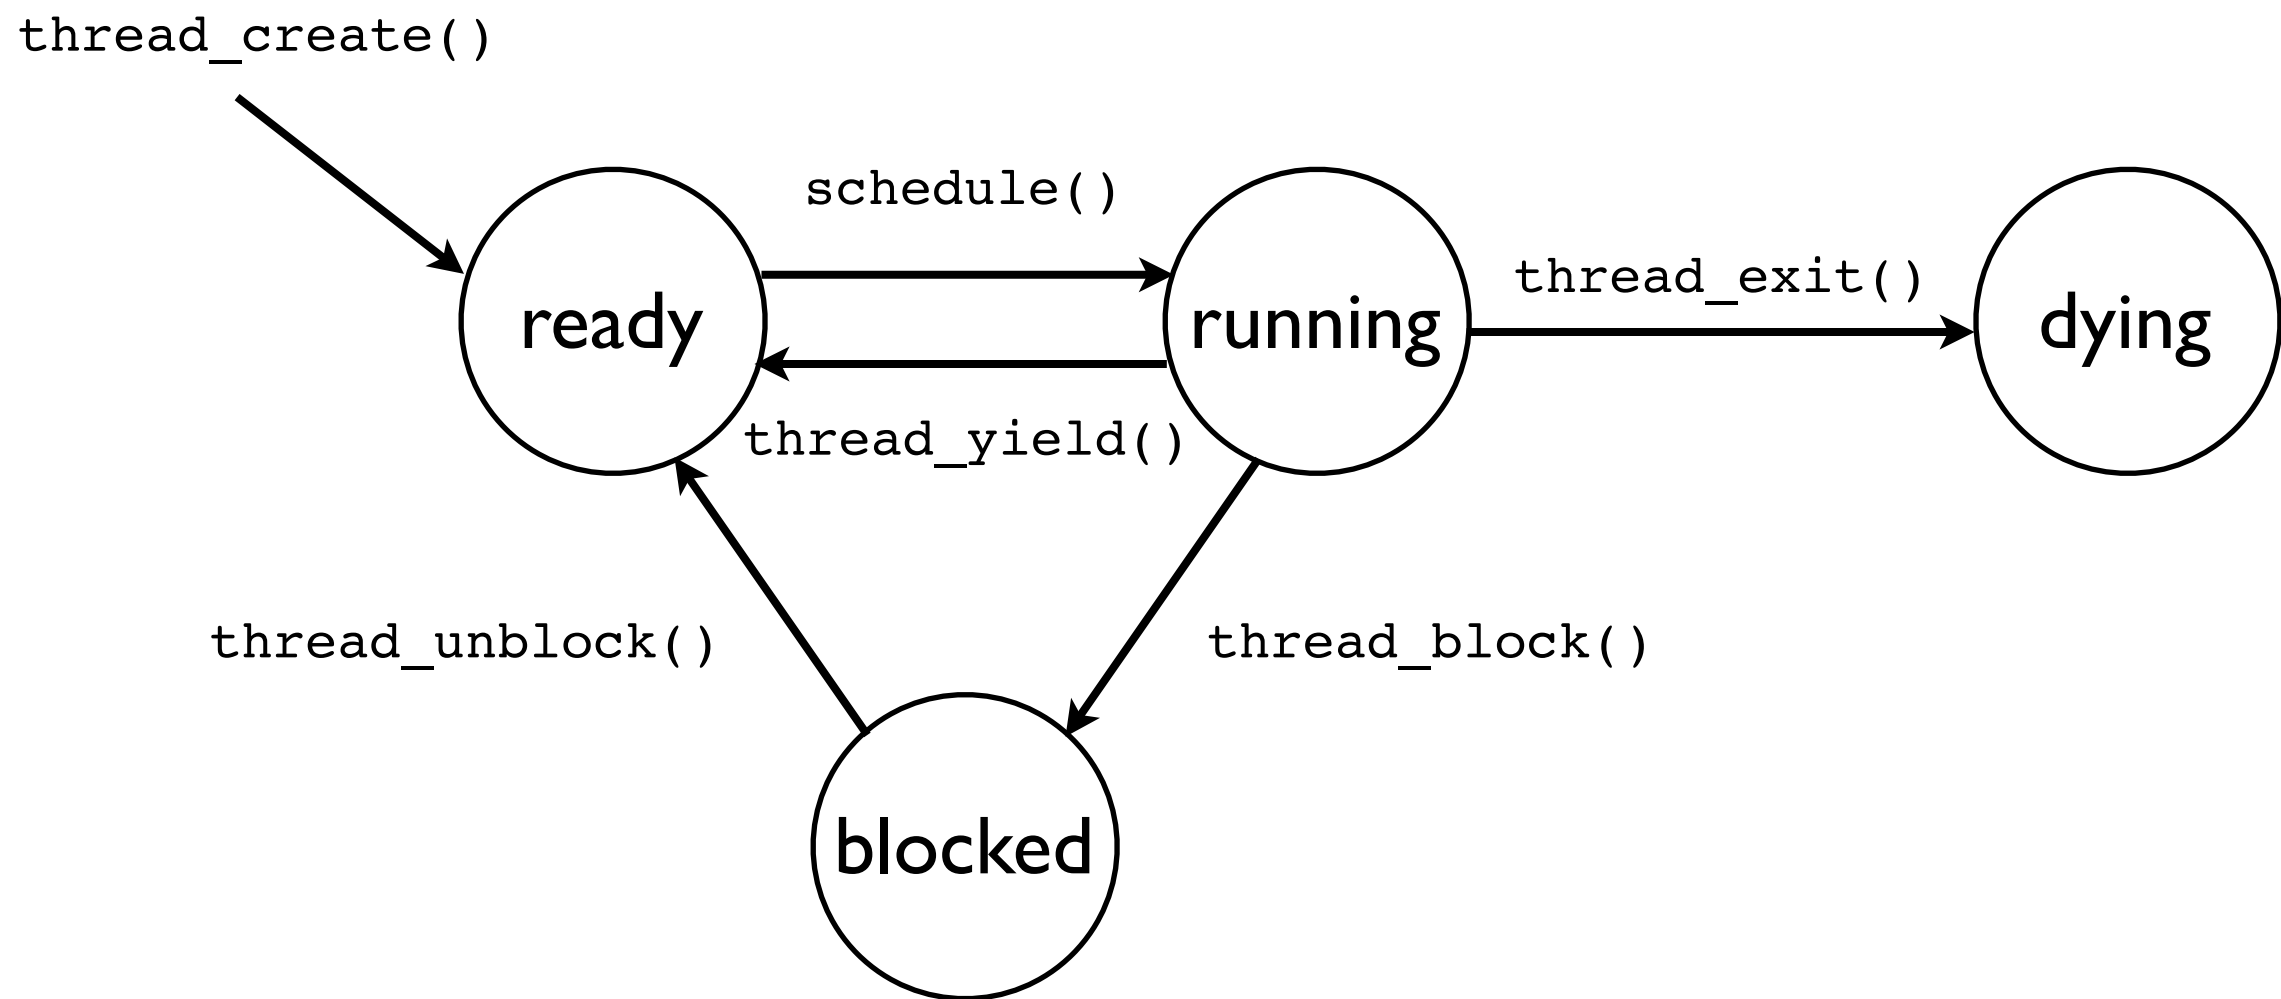
- Kernel maintains Thread Control Blocks
(see `struct thread` in `threads.h`)
 - `tid` (thread id)
 - thread name (for debugging)
 - stack pointer
 - status (ready, running, blocked, dying)

Threads

- Each thread gets a page of memory




Thread states



Scheduling

- Preemptive round-robin scheduler
 - Executes one thread at a time
(only one thread is in the running state)
 - Each thread is given a time slice
 - Switch threads when time slice expires

Ready list



- Push current thread to tail of the list
- Head of the list is the next thread to run

Scheduler

```
static void
schedule (void)
{
    struct thread *cur = running_thread ();
    struct thread *next = next_thread_to_run ();
    ...
    if (cur != next)
        prev = switch_threads (cur, next);
    ...
}

static struct thread *
next_thread_to_run (void)
{
    if (list_empty (&ready_list))
        return idle_thread;
    else
        return list_entry (list_pop_front (&ready_list), struct thread, elem);
}
```

Thread Switching

- The magic happens in `switch_threads(cur, next)`
 - Save registers on the stack
 - Save CPU's stack pointer in `cur->stack`
 - CPU stack pointer gets `next->stack`
 - Restore registers from stack

Timer interrupts

- Timer interrupts (see `devices/timer.c`)
 - Generates `TIMER_FREQ` “ticks” per second
 - Threads preempted every `TIME_SLICE` ticks

Timer interrupts

```
/* timer.c */
static void
timer_interrupt (struct intr_frame *args UNUSED)
{
    ticks++;
    thread_tick ();
}
```

```
/* thread.c */
void
thread_tick (void)
{
    struct thread *t = thread_current ();
    /* Update statistics. */
    ...
    /* Enforce preemption. */
    if (++thread_ticks >= TIME_SLICE)
        intr_yield_on_return ();
}
```

Timer interrupts

```
/* timer.c */
static void
timer_interrupt (struct intr_frame *args UNUSED)
{
    ticks++;
    thread_tick ();
}
```

```
/* thread.c */
void
thread_tick (void)
{
    struct thread *t = thread_current ();
    /* Update statistics. */
    ...
    /* Enforce preemption. */
    if (++thread_ticks >= TIME_SLICE)
        intr_yield_on_return ();
}
```

`intr_yield_on_return()`
causes the interrupt
handler to call
`thread_yield()` before
returning

Synchronization

- Pintos provides several synchronization primitives
 - Semaphores
 - Locks
 - Condition variables
- We will study these mechanisms later on...

Synchronization

- For now: access to shared data should be protected by disabling interrupts

```
enum intr_level old_level;  
old_level = intr_disable ();  
  
/* critical section */  
  
intr_set_level (old_level);
```

Assignment

- Reimplement `timer_sleep()`
- Current implementation

```
void
timer_sleep (int64_t ticks)
{
    int64_t start = timer_ticks ();
    while (timer_elapsed (start) < ticks)
        thread_yield ();
}
```

- What is the problem???

Assignment

- Reimplement `timer_sleep()`
- Current implementation

```
void
timer_sleep (int64_t ticks)
{
    int64_t start = timer_ticks ();
    while (timer_elapsed (start) < ticks)
        thread_yield ();
}
```

- What is the problem??? Busy wait!

Assignment

- Your submission should
 - avoid busy waiting
 - not brake tests that already pass!
 - ideally, be efficient

Assignment

\$ make check

....

pass tests/threads/alarm-single

pass tests/threads/alarm-multiple

pass tests/threads/alarm-simultaneous

FAIL tests/threads/alarm-priority

pass tests/threads/alarm-zero

pass tests/threads/alarm-negative

FAIL tests/threads/priority-change

FAIL tests/threads/priority-donate-one

FAIL tests/threads/priority-donate-multiple

FAIL tests/threads/priority-donate-multiple2

FAIL tests/threads/priority-donate-nest

FAIL tests/threads/priority-donate-sema

FAIL tests/threads/priority-donate-lower

FAIL tests/threads/priority-fifo

FAIL tests/threads/priority-preempt

FAIL tests/threads/priority-sema

FAIL tests/threads/priority-condvar

FAIL tests/threads/priority-donate-chain

FAIL tests/threads/mlfqs-load-1

FAIL tests/threads/mlfqs-load-60

FAIL tests/threads/mlfqs-load-avg

FAIL tests/threads/mlfqs-recent-1

pass tests/threads/mlfqs-fair-2

pass tests/threads/mlfqs-fair-20

FAIL tests/threads/mlfqs-nice-2

FAIL tests/threads/mlfqs-nice-10

FAIL tests/threads/mlfqs-block

20 of 27 tests failed.

Hints

- Use `thread_block()` and keep track of sleeping threads
- During timer interrupts, `thread_unblock()` sleeping threads if their timer expired
- Disable interrupts to protect critical sections

Readings

- Read pintos docs!
- Chapter 2 up to 2.2.2
- Appendix A.2
- Appendix A.3.1
- Skim through Appendix A.1