

Pintos User Programs II

Executing child processes
Controlling memory access

Recap

- So far, we know how to:
 - run and debug pintos tests
 - add our own tests
 - use `lib/kernel/list.h`
 - implement `sleep()` with no busy-wait, thread priority and niceness and system calls to be used by user programs

Next steps

- Today we implement two syscalls:
`wait()` and `exec()`
- To implement `exec()`, we need:
 - synchronization
(with `thread_block()`, or just a semaphore)
 - check the pointers passed by the user
(requires some understanding of memory paging)

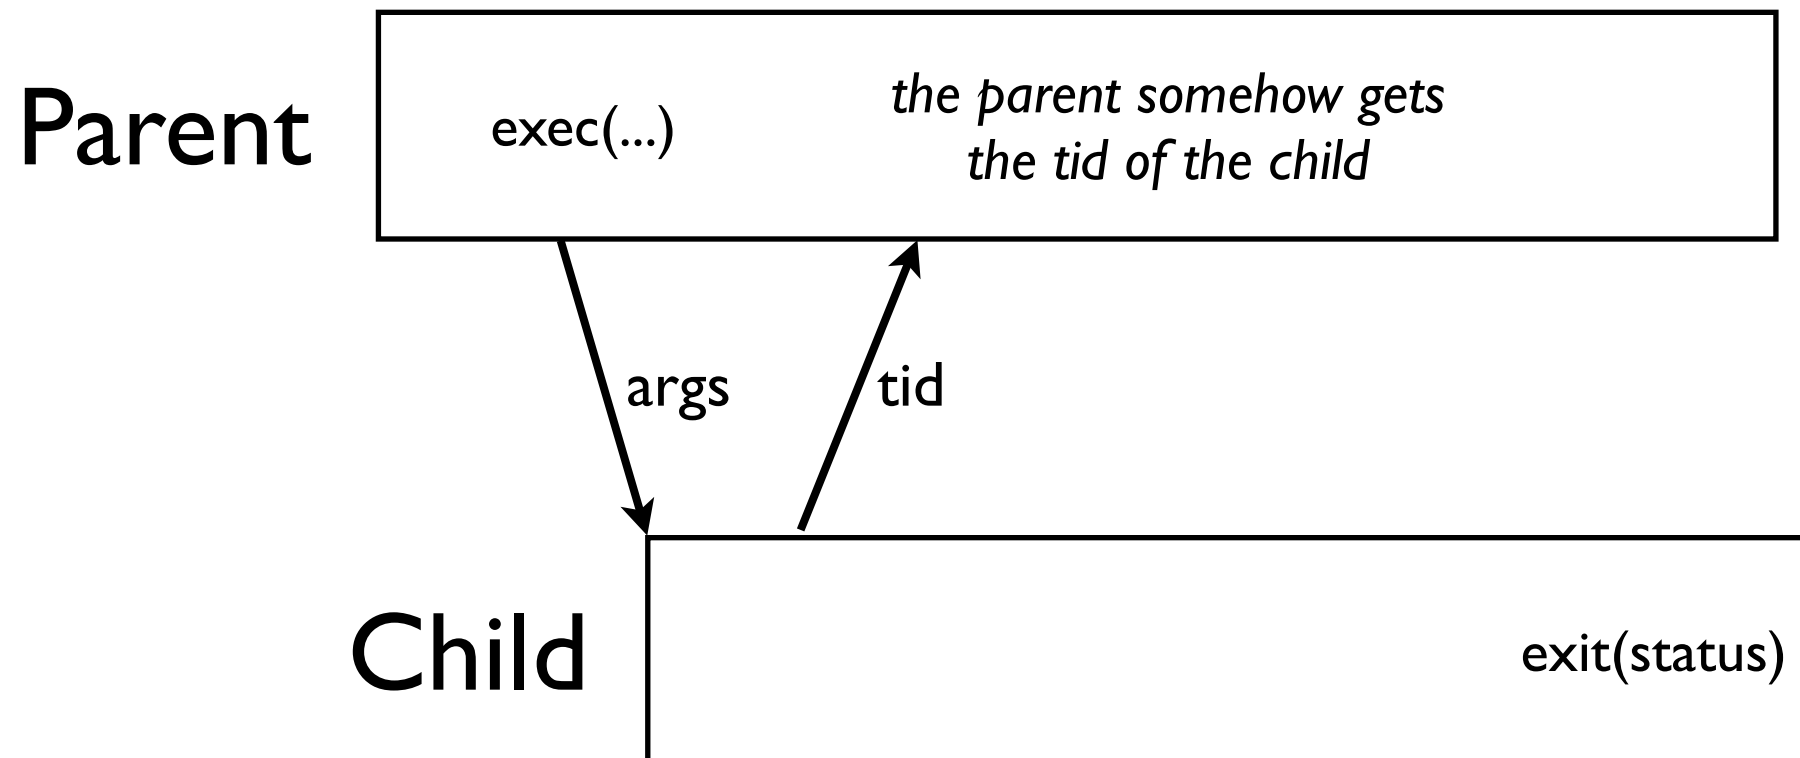
The `wait ()` system call

- `int wait (pid_t pid)`
- The calling process blocks until its child `p`, which has id `pid`, has finished
- The return value must be:
 - the exit status of `p`, if all was fine
(even if `p` finished before `wait` was called!)
 - `-1` if `p` was killed
 - `-1` if `p` is not a direct child of the caller
(this makes it easier)
 - `-1` if the calling process already waited for that `pid` before

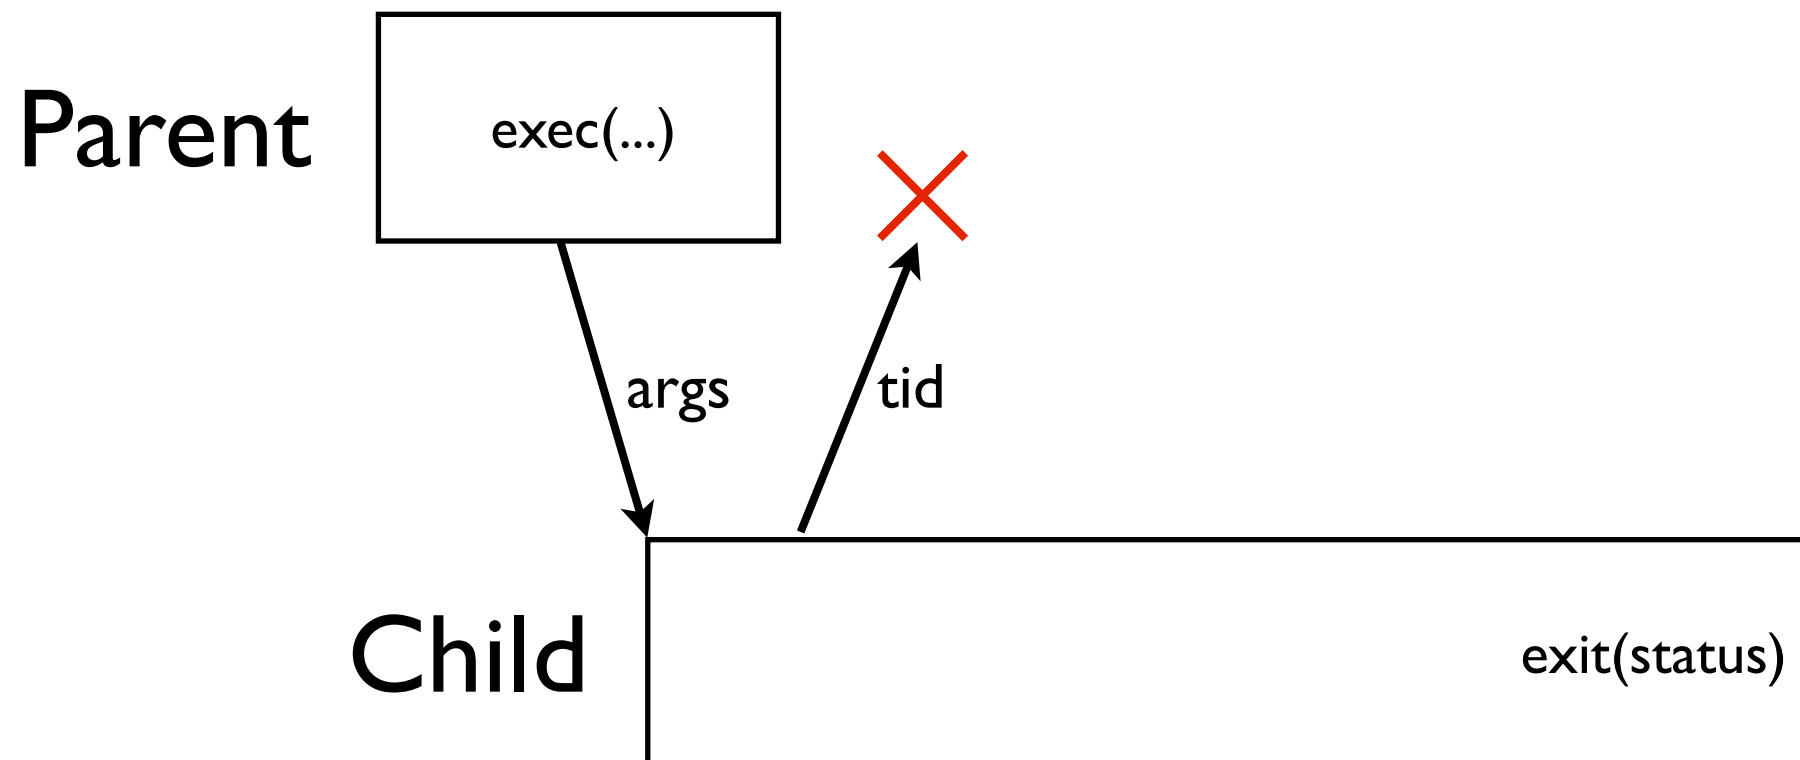
The `exec ()` system call

- `pid_t exec (const char *cmd_line)`
- A user program can use it to execute a command, creating a child process
- It tries to run the executable given in `cmd_line`
- The return value is the pid (tid) of the child process (or -1 in case of error)

Executing - `exec(...)`

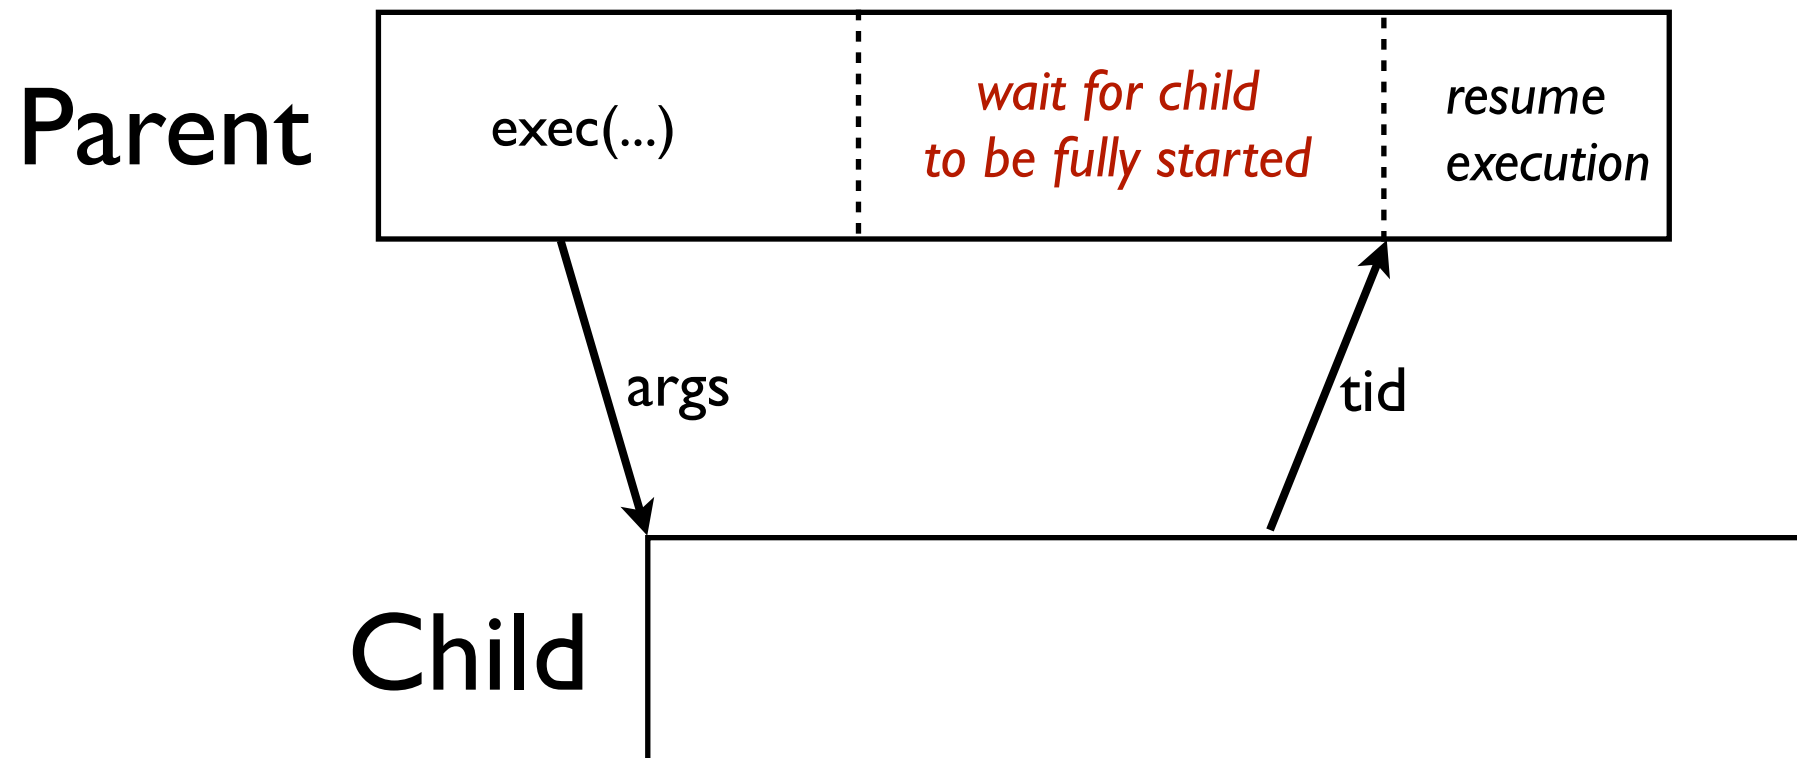


Executing - `exec(...)`



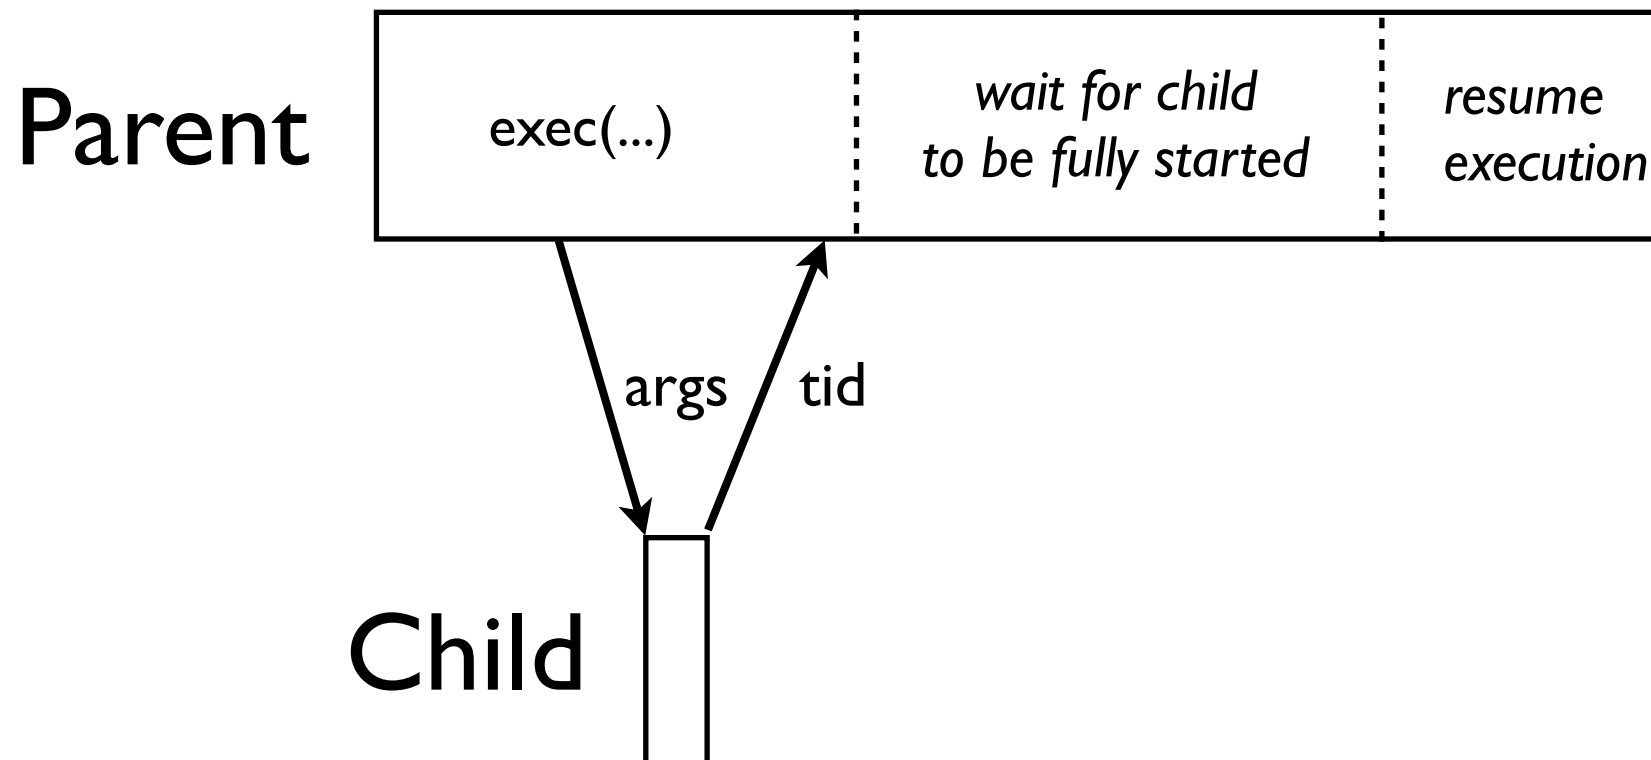
The parent thread must wait for the child creation

Executing - `exec(...)`



The parent thread must wait for the child creation

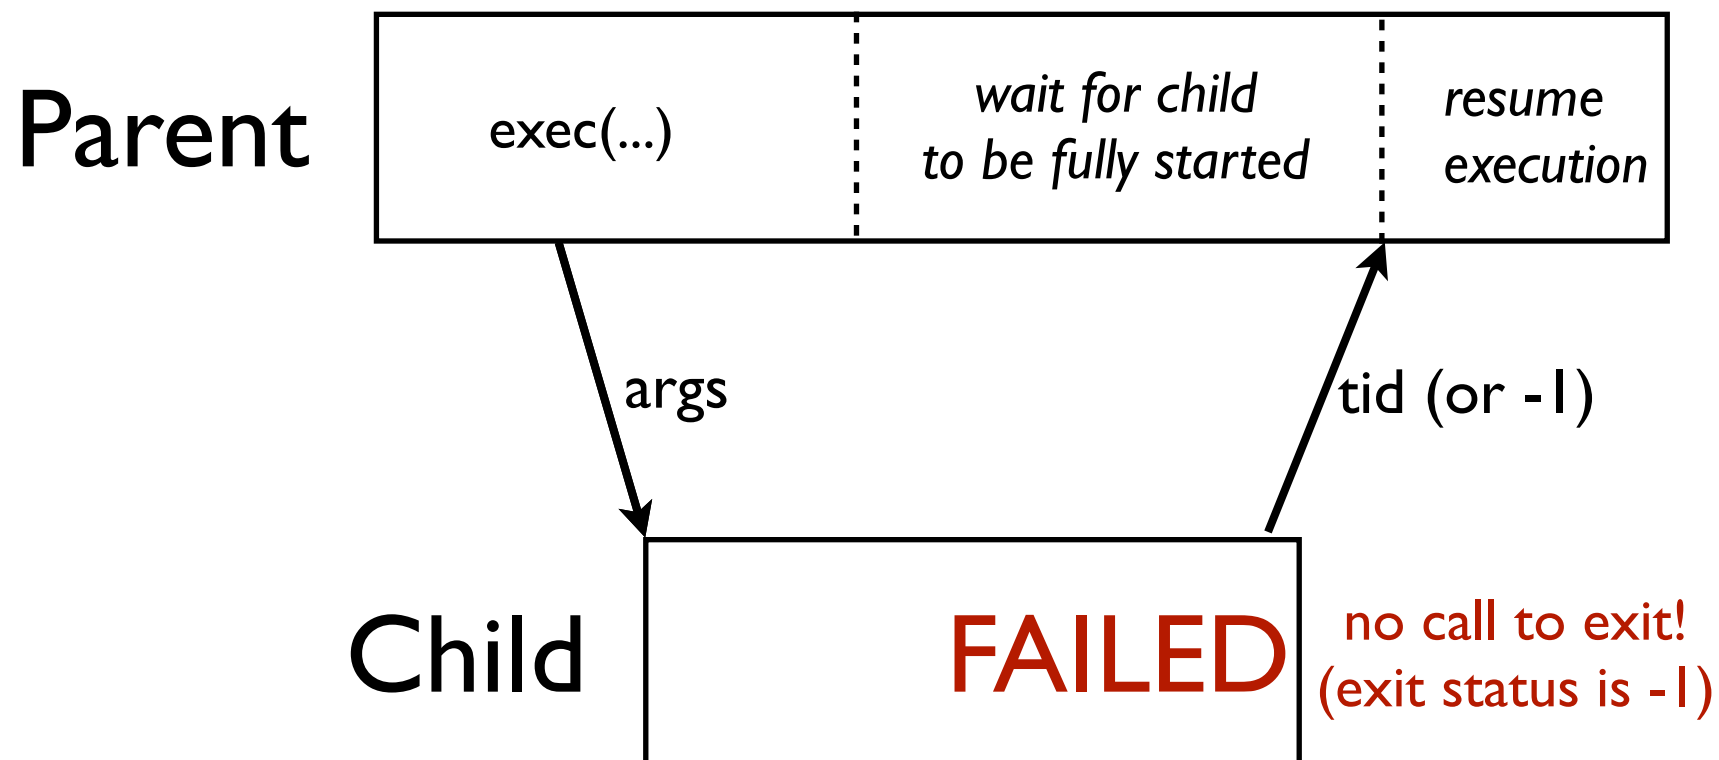
Executing - `exec(...)`



The parent thread must wait for the child creation

The child thread may finish too fast

Executing - `exec(...)`



The parent thread must wait for the child creation
The child thread may finish too fast
The child thread may fail before telling some status

Pintos semaphores

- Used for synchronizing threads
- Can solve the (wait for child) problem of the `exec syscall`
- Available including “thread/synch.h”
- Provided API:

```
struct semaphore s;  
sema_init(&s, 0);  
sema_up(&s);  
sema_down(&s);
```

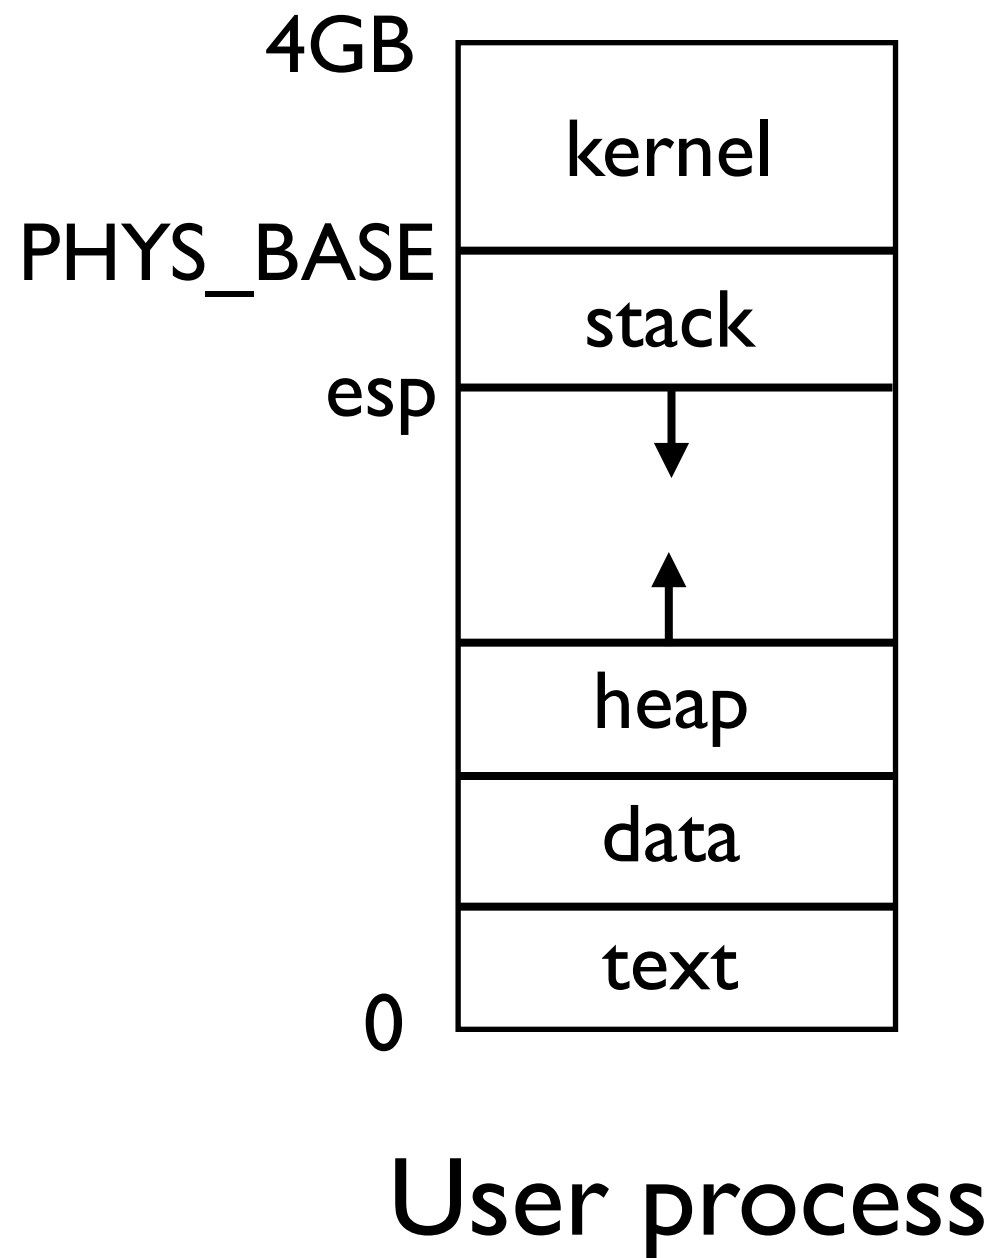
Pintos semaphores API

- `struct semaphore`
semaphore type; must be initialized with `sema_init`
- `sema_init (struct semaphore * s, unsigned val)`
initialize semaphore pointed by `s` with value `val`
- `sema_down (struct semaphore * s)`
if `s` is 0 (zero), block the calling thread and put it in a list waiting for `s`;
otherwise, decrement `s` by 1
- `sema_up (struct semaphore * s)`
if `s` is 0 (zero) and there is some thread waiting for `s`, unblock one of the
threads that are waiting for `s`; otherwise, increment `s` by 1
- neither `sema_up` or `sema_down` can be interrupted (they're *atomic*)

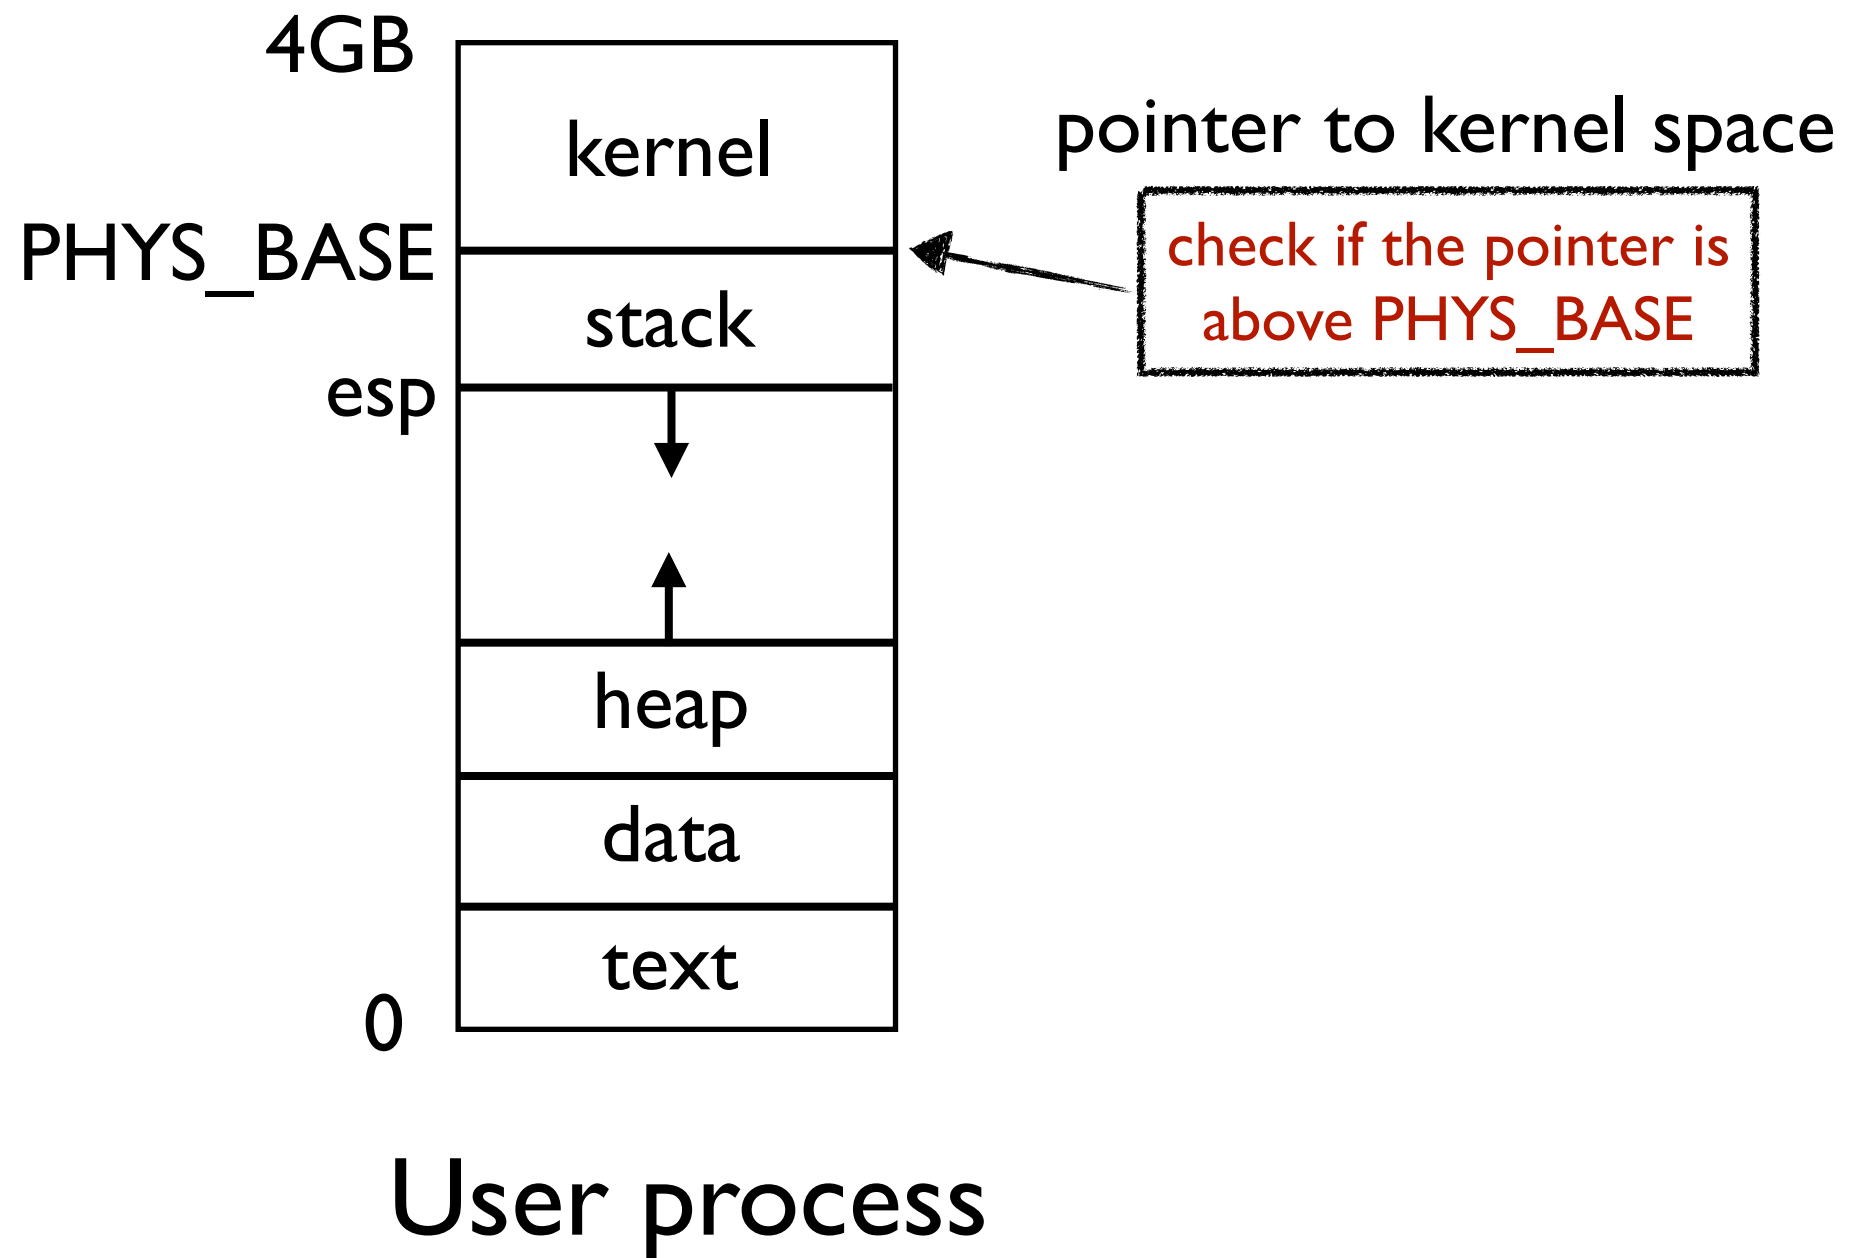
Memory access

- The user may provide an invalid pointer in a syscall
 - a null pointer
 - a pointer to kernel address space
 - a pointer to unmapped virtual memory
- The kernel (you) should control this

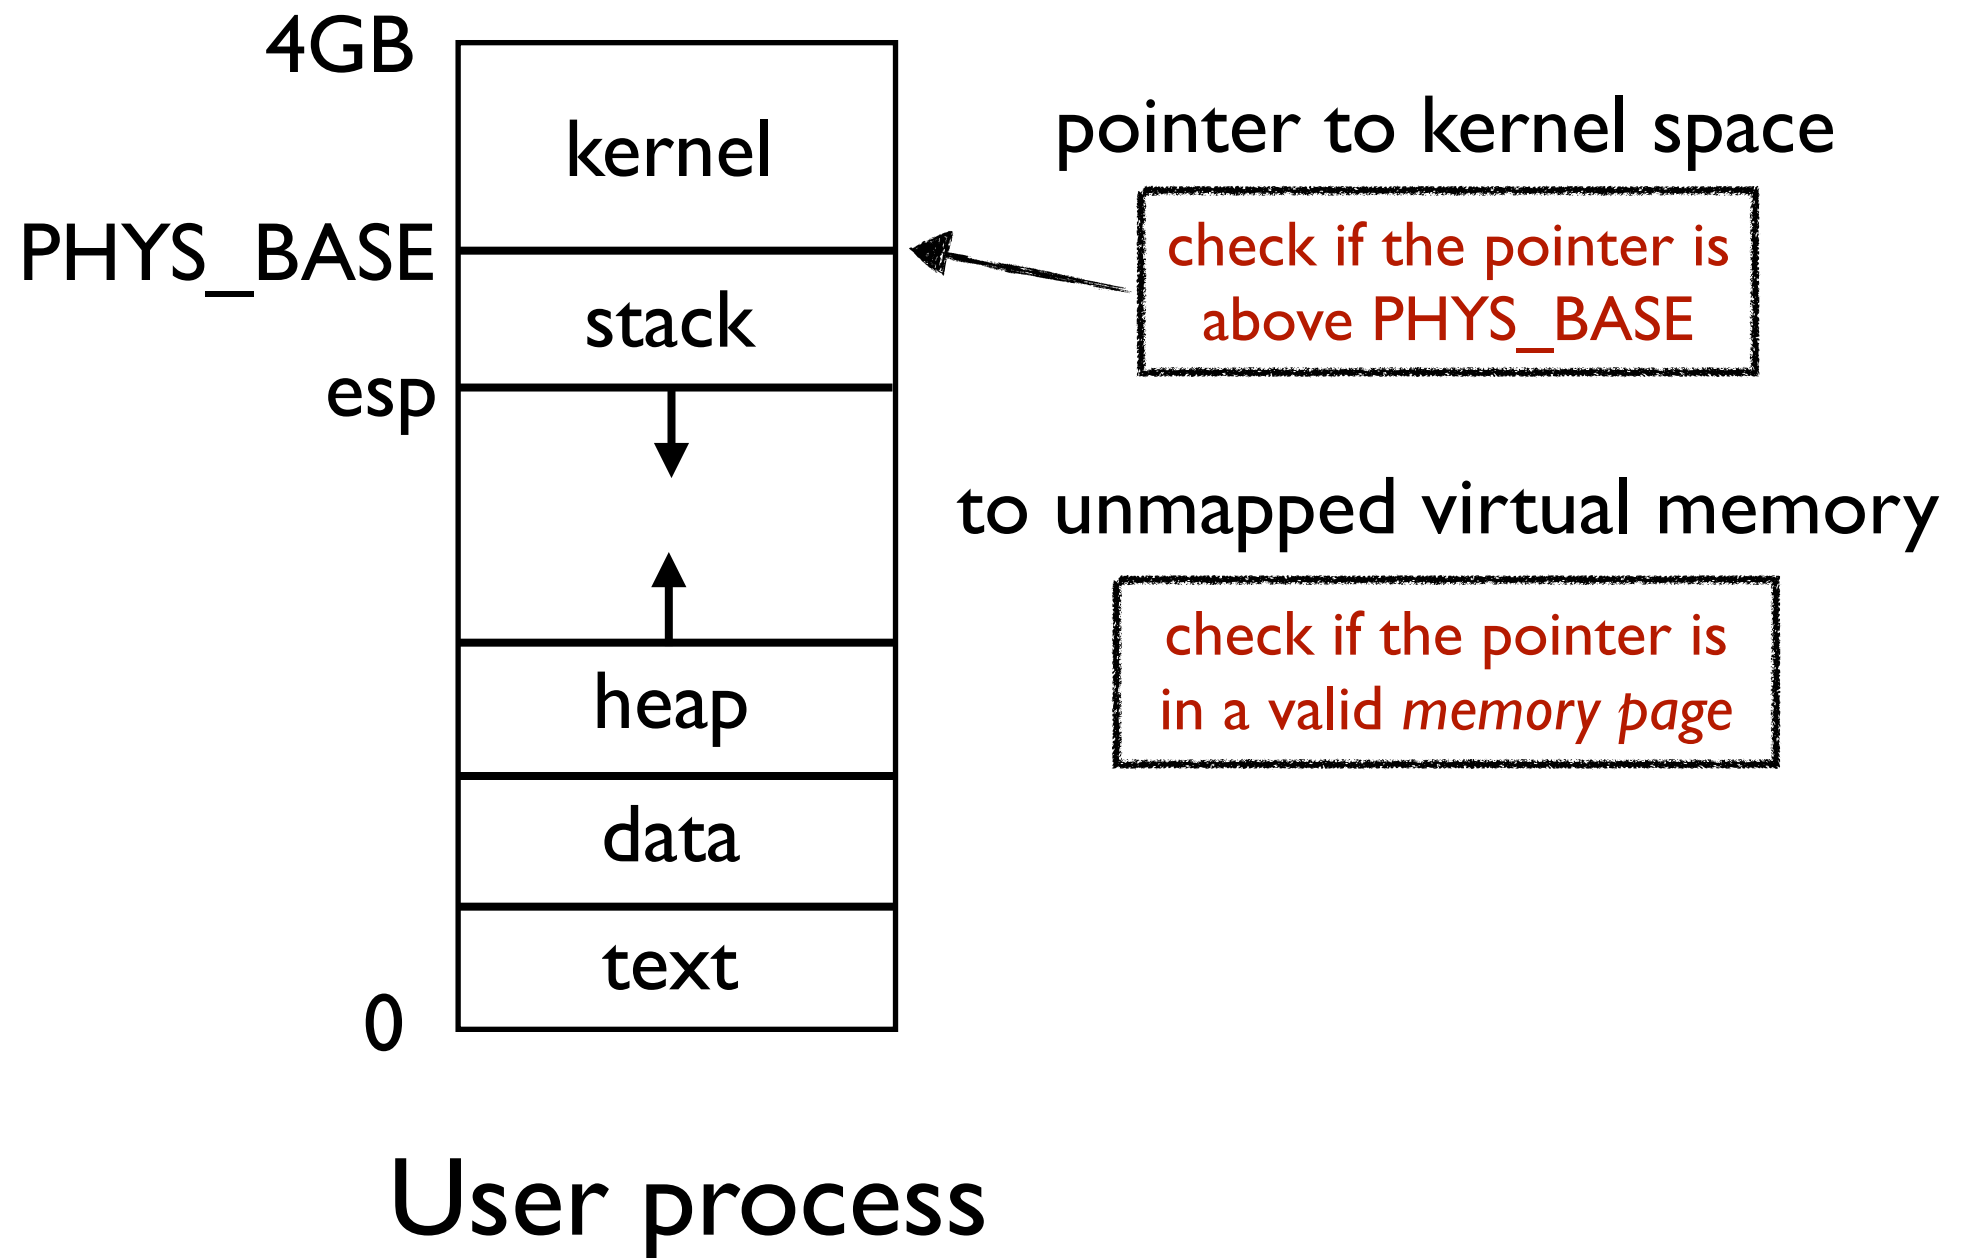
Memory layout




Memory layout



Memory layout



Memory paging



		Page
0	a	0
1	b	
2	c	
3	d	
4	e	1
5	f	
6	g	
7	h	
8	i	2
9	j	
10	k	
11	l	
12	m	3
13	n	
14	o	
15	p	

logical memory
Pintos: virtual address

0	5
1	6
2	1
3	2


page table
Pintos: page directory
(thread->pagedir)

0		Frame
		0
4	i j k l	1
8	m n o p	2
12		3
16		4
20	a b c d	5
24	e f g h	6
28		7

physical memory

The process allocates memory on demand
Still, though, its virtual memory looks contiguous

Memory paging



Page	
0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

logical memory
Pintos: virtual address

0	5
1	6
2	1
3	2

page table
Pintos: page directory
(struct thread.pagedir)

check if the user pointer is in
a valid memory page:
hints in [vaddr.h](#) and [pagedir.c](#)

Frame	
0	
4	i j k l
8	m n o p
12	
16	
20	a b c d
24	e f g h
28	

physical memory

The process allocates memory on demand
Still, though, its virtual memory looks contiguous

Tests

- After implementing, you should pass:
 - exec-once
 - exec-arg
 - exec-multiple
 - exec-missing
 - exec-bad-ptr
 - wait-simple
 - wait-twice
 - wait-bad-pid

Readings

- Chapter 3:
specially sections 3.1.4 and 3.3.4